

SII-RLVE: Shaping Inter-task Imbalance with Gradient-equilibration for RLVE

Ming Wen
253108090054


Shixiang Song
253108120163

Muze Wang
253108120135



Motivation: The Leaky Bucket

Train Observation


Adaptive accuracy per env (paper-16 pool)






Eval Observation



SII-RLVE



Empirical Results



Model & Metric Configuration	Best Eval Score
Qwen-2.5-7B rlve	0.3380
Qwen-2.5-7B sii-rlve	0.3632
DeepSeek-distill-qwen-7B rlve	0.4110
DeepSeek-distill-qwen-7B sii-rlve	0.4305

Table 1. Best Evaluation Results

Fig 2. Comparison of training accuracy of 8 envs

Fig 3. Comparison of sampling rate versus steps